

1 Methodology

As discussed in Section 1, our approach leverages deep reinforcement learning¹ to align outputs with human preferences.

1.1 Data Processing

The corpus is tokenized using Byte-Pair Encoding (BPE) with a vocabulary size of 50,257.

“Alignment remains one of the most critical open problems in modern artificial intelligence, bridging the gap between raw computational capability and safe, intended utility.”

¹Specifically, Proximal Policy Optimization (PPO).